

1. Write a program for the following

a.To generate an array of random numbers from a normal distribution for the array of a given shape.

```
import numpy as np
# Enter the value of n
n=int(input('Enter no. of values:'))
# Generates n random numbers from Normal Distribution
rand_num = np.random.normal(0,1,n)
print(n, " random numbers from a standard normal distribution:")
print(rand_num)
arr=np.array([rand_num])
# Displays the size of the Array
print(arr.shape)
```

output

```
Enter no. of values:10
10 random numbers from a standard normal distribution:
[-0.34953998  1.60514591 -0.60005696  0.26263808  0.87930153
 0.98339437
 0.40472381 -0.73362668 -0.20067116 -0.97191095]
(1, 10)
```

b. Implement Arithmetic operations on two arrays (perform broadcasting also.)

```
#Generates an Array A with 0 to 11 in a 3X4 form
A = np.arange(12).reshape(3,4)
print(A)
#Generates an Array B with 0 to 3 in a 1X3 form
B = np.arange(4)
print(B)
#Performs the addition between A and B
c=A+B
print(c)
#Similarly perform remaining arithmetic operations (Subtraction,multiplication, division )
```

output

```
[[ 0  1  2  3]
 [ 4  5  6  7]
 [ 8  9 10 11]]
[0 1 2 3]
[[ 0  2  4  6]
 [ 4  6  8 10]
 [ 8 10 12 14]]
```

c. Find minimum, maximum, mean in a given array. (in both the axes)

```
arr = np.array([[11, 2, 3],[4, 5, 16],[7, 81, 22]])
# finding the maximum and minimum element in the array
max_element = np.max(arr)
min_element = np.min(arr)
```

```
# printing the result
```

```

print('maximum element in the array is:', max_element)
print('minimum element in the array is:', min_element)
# finding the maximum and
# minimum element in the array
max_element_column = np.max(arr, 0)
max_element_row = np.max(arr, 1)

min_element_column = np.amin(arr, 0)
min_element_row = np.amin(arr, 1)

# printing the result
print('maximum elements in the columns of the array is:', max_element_column)

print('maximum elements in the rows of the array is:', max_element_row)

print('minimum elements in the columns of the array is:', min_element_column)

print('minimum elements in the rows of the array is:', min_element_row)

# mean of the flattened array
print("\nmean of arr, axis = None : ", np.mean(arr))

# mean along the axis = 0 (row-wise)
print("\nmean of arr, axis = 0 : ", np.mean(arr, axis = 0))

# mean along the axis = 1 (Column-wise)
print("\nmean of arr, axis = 1 : ", np.mean(arr, axis = 1))

```

output

```

maximum element in the array is: 81
minimum element in the array is: 2
maximum elements in the columns of the array is: [11 81 22]
maximum elements in the rows of the array is: [11 16 81]
minimum elements in the columns of the array is: [4 2 3]
minimum elements in the rows of the array is: [2 4 7]

mean of arr, axis = None : 16.777777777777778

mean of arr, axis = 0 : [ 7.33333333 29.33333333 13.66666667]

mean of arr, axis = 1 : [ 5.33333333  8.33333333 36.66666667]

```

d. Implement np.arange and np.linspace functions.

```

# Prints all numbers from 0 to 9 in steps of 1
arr=np.linspace(start = 0, stop = 10, num = 11, dtype = int)
print(arr)
arr=np.linspace(start = 0, stop = 1, num = 11)
print(arr)
# Prints all numbers from 1 to 2 in steps of 0.1

```

```
print(np.arange(1, 2, 0.1))
```

output

```
[ 0  1  2  3  4  5  6  7  8  9 10]
[0.  0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1. ]
[1.  1.1 1.2 1.3 1.4 1.5 1.6 1.7 1.8 1.9]
```

e. Create a pandas series from a given list.

```
# import pandas lib. as pd
import pandas as pd
# Assume l1 is a list of the following words
l1 = ['ZERO', 'ONE', 'TWO', 'THREE',
      'FOUR', 'FIVE', 'SIX', 'SEVEN', 'EIGHT', 'NINE', 'TEN']
```

```
# create Pandas Series with define indexes
x = pd.Series(l1)
```

```
# print the Series
print(x)
```

output

```
0      ZERO
1       ONE
2       TWO
3    THREE
4     FOUR
5     FIVE
6      SIX
7    SEVEN
8    EIGHT
9     NINE
10     TEN
dtype: object
```

f. Create pandas series with data and index and display the index values.

```
# import pandas lib. as pd
import pandas as pd

# create Pandas Series with define indexes
x = pd.Series([10, 20, 30, 40, 50], index=['a', 'b', 'c', 'd', 'e'])
```

```
# print the Series
print(x)
```

output

```
a      10
b      20
c      30
```

```
d      40
e      50
dtype: int64
```

- g. Create a data frame with columns at least 5 observations**
- i. select a particular column from the DataFrame**
- ii. Summarize the data frame and observe the stats of the DataFrame created**
- iii. Observe the mean and standard deviation of the data frame and print the values.**

```
import pandas as pd
import numpy as np
```

```
exam_data = {'name': ['Anastasia', 'Dima', 'Katherine', 'James', 'Emily'], 'score': [12.5, 9, 16.5,
np.nan, 9], 'attempts': [1, 3, 2, 3, 2],
             'qualify': ['yes', 'no', 'yes', 'no', 'no']}
labels = ['a', 'b', 'c', 'd', 'e']
```

```
df = pd.DataFrame(exam_data , index=labels)
print("Dataset is as follows")
print(df)
print("Summary of the Dataset")
print(df.info())
print("Statistical values of numerical attributes")
print(df.describe())
meanvalue=df.score.mean()
stdvalue=df.score.std()
print('mean value of Score is',meanvalue)
print('Standard deviation of score is',stdvalue)
```

output

```
Dataset is as follows
   name  score  attempts  qualify
a  Anastasia   12.5         1     yes
b      Dima     9.0         3     no
c  Katherine   16.5         2     yes
d      James    NaN         3     no
e      Emily     9.0         2     no
Summary of the Dataset
<class 'pandas.core.frame.DataFrame'>
Index: 5 entries, a to e
Data columns (total 4 columns):
 #   Column        Non-Null Count  Dtype
---  -
 0   name          5 non-null      object
 1   score         4 non-null      float64
 2   attempts     5 non-null      int64
 3   qualify       5 non-null      object
dtypes: float64(1), int64(1), object(2)
memory usage: 200.0+ bytes
None
Statistical values of numerical attributes
           score  attempts
count    4.000000    5.00000
mean     11.750000    2.20000
```

```

std      3.570714    0.83666
min      9.000000    1.00000
25%      9.000000    2.00000
50%     10.750000    2.00000
75%     13.500000    3.00000
max     16.500000    3.00000
mean value of Score is 11.75
Standard deviation of score is 3.570714214271425

```

1. Write a Program to determine the following in the Titanic Survival [data](#).

a. Determine the data type of each column.

```

# importing all the necessary libraries
import pandas as pd
import numpy as np

#we need to read the data
data = pd.read_csv("https://raw.githubusercontent.com/naveenjoshii/Intro-to-MachineLearning/master/Titanic/titanic.csv")
#print top 5 rows
print(data.head())

```

output

PassengerId	Survived	Pclass	...	Fare	Cabin	Embarked
0	1	0	3 ...	7.2500	NaN	S
1	2	1	1 ...	71.2833	C85	C
2	3	1	3 ...	7.9250	NaN	S
3	4	1	1 ...	53.1000	C123	S
4	5	0	3 ...	8.0500	NaN	S

[5 rows x 12 columns]

```

# to get the datatype of all columns we can use Dataframe.dtypes
print(data.dtypes)

```

output

```

PassengerId      int64
Survived          int64
Pclass            int64
Name              object
Sex               object
Age              float64
SibSp             int64
Parch             int64
Ticket            object
Fare              float64
Cabin             object
Embarked          object
dtype: object

```

b. Find the number of non-null values in each column.

```
# Dataframe.info() gives all information about every column in our dataset
data.info()
```

output

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   PassengerId      891 non-null    int64  
 1   Survived         891 non-null    int64  
 2   Pclass          891 non-null    int64  
 3   Name             891 non-null    object  
 4   Sex              891 non-null    object  
 5   Age             714 non-null    float64 
 6   SibSp           891 non-null    int64  
 7   Parch           891 non-null    int64  
 8   Ticket           891 non-null    object  
 9   Fare            891 non-null    float64 
10   Cabin           204 non-null    object  
11   Embarked         889 non-null    object  
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

c. Find out the unique values in each categorical column and frequency of each unique value.

```
# categorical is nothing but the datatype which is other than numerical datatype (i.e int,float etc).
```

```
# to get the all categorical columns, we can use Dataframe.select_dtypes and we have to specify which
```

```
#datatype we required.
```

```
# In our case it would be "object" datatype
```

```
categorical_cols = data.select_dtypes(include=['object']).columns.tolist()
```

```
print("Categorical columns are : ",categorical_cols)
```

```
print("printing the results")
```

```
for i in categorical_cols:
```

```
    print("===== Column '"+i+"' =====")
```

```
    print(data[i].value_counts())
```

output

```
Categorical columns are :  ['Name', 'Sex', 'Ticket', 'Cabin',
'Embarked']
printing the results
===== Column 'Name' =====
Robert, Mrs. Edward Scott (Elisabeth Walton McMillan)    1
Smith, Mr. Thomas                                          1
Cameron, Miss. Clear Annie                                1
Parkes, Mr. Francis "Frank"                               1
```

```

Panula, Mrs. Juha (Maria Emilia Ojala) 1
Walker, Mr. William Anderson 1
Hassab, Mr. Hammad 1
Olsen, Mr. Karl Siegwart Andreas 1
Reed, Mr. James George 1
Wiseman, Mr. Phillippe 1
Name: Name, Length: 891, dtype: int64
===== Column 'Sex' =====
male 577
female 314
Name: Sex, dtype: int64
===== Column 'Ticket' =====
1601 7
CA. 2343 7
347082 7
3101295 6
CA 2144 6
350034 1
19947 1
A/5 21174 1
PC 17474 1
SOTON/OQ 392082 1
Name: Ticket, Length: 681, dtype: int64
===== Column 'Cabin' =====
G6 4
B96 B98 4
C23 C25 C27 4
F2 3
E101 3
B38 1
B102 1
E58 1
C101 1
B4 1
Name: Cabin, Length: 147, dtype: int64
===== Column 'Embarked' =====
S 644
C 168
Q 77
Name: Embarked, dtype: int64

```

d. Find the number of rows where age is greater than the mean age of data.

```

# to get mean of age column
age_mean = data['Age'].mean()
print("Mean of Age is : ",age_mean)
print("printing the result")
print(np.sum(data['Age']>age_mean))

```

output

```

Mean of Age is : 29.69911764705882
printing the result
330

```

e. Delete all the rows with missing values.

```
print("length of dataframe before deleting rows with missing values",len(data))
# deletes the rows where at least one element is missing
data.dropna(inplace=True)
print("length of dataframe after the deletion of missing value rows",len(data))
```

output

```
length of dataframe before deleting rows with missing values 891
length of dataframe after the deletion of missing value rows 183
```

3. Perform Data Analysis on the Titanic Data Set to answer the following.

```
#importing all the necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
#reading data
data = pd.read_csv("https://raw.githubusercontent.com/naveenjoshii/Intro-to-MachineLearning/master/Titanic/titanic.csv")
print(data.head())
```

output

	PassengerId	Survived	Pclass	...	Fare	Cabin	Embarked
0	1	0	3	...	7.2500	NaN	S
1	2	1	1	...	71.2833	C85	C
2	3	1	3	...	7.9250	NaN	S
3	4	1	1	...	53.1000	C123	S
4	5	0	3	...	8.0500	NaN	S

```
[5 rows x 12 columns]
```

a. Information regarding each column of the data

```
#printing the info about all the columns
print(data.info())
```

output

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   PassengerId     891 non-null   int64
```



```

1   Survived      891 non-null    int64
2   Pclass        891 non-null    int64
3   Name          891 non-null    object
4   Sex           891 non-null    object
5   Age           714 non-null    float64
6   SibSp         891 non-null    int64
7   Parch         891 non-null    int64
8   Ticket        891 non-null    object
9   Fare          891 non-null    float64
10  Cabin         204 non-null    object
11  Embarked      889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
None

```

b. Impact of each column on the label

plotting the correlation using heatmap


```

sns.heatmap(data.corr(),cmap='coolwarm',xticklabels=True,annot=True)
plt.title('data.corr()')

```

output

```
Text(0.5, 1.0, 'data.corr()')
```



c. Number of survivals in each gender


plotting countplot for Each gender who has survived and not survived

```
sns.set_style('whitegrid')
```

```
sns.countplot(x='Survived',hue='Sex',data=data,palette='colorblind')
```

output

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f621a047810>
```




d. Number of survivals in each passenger class

```
#plotting count plot for no of survivals in each class
sns.set_style('whitegrid')
sns.countplot(x='Survived',hue='Pclass',data=data,palette='bright')
```

output

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f621a034510>
```




e. The number of people who are not alone.

```
# count plot for who has siblings/spouse
sns.countplot(x = 'SibSp', data = data,)
```

output

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6219b36390>
```



4. Perform Data Analysis on the California House Price data to answer the following

```
# importing all the necessary libraries
import pandas as pd
import numpy as np
```

```
#we need to read the data
data = pd.read_csv("https://raw.githubusercontent.com/ageron/handson-ml/master/datasets/housing/housing.csv")
#print top 5 rows
print(data.head())
```

output

```
longitude  latitude  ...  median_house_value  ocean_proximity
0    -122.23    37.88  ...           452600.0      NEAR BAY
1    -122.22    37.86  ...           358500.0      NEAR BAY
2    -122.24    37.85  ...           352100.0      NEAR BAY
3    -122.25    37.85  ...           341300.0      NEAR BAY
4    -122.25    37.85  ...           342200.0      NEAR BAY
```

```
[5 rows x 10 columns]
```

a. Data Type of each column and info regarding each column

```
# data information for each column
print(data.info())
```

Output

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
 #   Column              Non-Null Count  Dtype
---  -
```

```

0    longitude      20640 non-null float64
1    latitude      20640 non-null float64
2    housing_median_age  20640 non-null float64
3    total_rooms    20640 non-null float64
4    total_bedrooms 20433 non-null float64
5    population     20640 non-null float64
6    households     20640 non-null float64
7    median_income  20640 non-null float64
8    median_house_value 20640 non-null float64
9    ocean_proximity 20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
None

```

b. The average age of a house in the data set.

```

# printing average age of house
print(data['housing_median_age'].mean())

```

Output

```

28.639486434108527

```

c. Determines top 10 localities with the high difference between income and house value. Also, top 10 localities that have the lowest difference

```

#calculating the difference btw House value and income and adding new column
'diff_income_and_house_value' with difference values
data['diff_income_and_house_value'] = data['median_house_value'] - data['median_income']
# sorting the whole dataframe by the difference value in descending order
data.sort_values(by='diff_income_and_house_value', ascending=False,inplace=True)
#printing the top 10 localities with highest difference
print("the top 10 localities with highest difference")
print(data['ocean_proximity'].head(10))
#printing the top 10 localities with lowest difference
print("the top 10 localities with lowest difference")
print(data['ocean_proximity'].tail(10))

```

Output

```

the top 10 localities with highest difference
4861    <1H OCEAN
6688    INLAND
16642    NEAR OCEAN
15661    NEAR BAY
15652    NEAR BAY
6639    <1H OCEAN
459     NEAR BAY
89      NEAR BAY
10448    <1H OCEAN
17819    <1H OCEAN
Name: ocean_proximity, dtype: object
the top 10 localities with lowest difference

```

```

2779          INLAND
16186         INLAND
14326    NEAR OCEAN
1825        NEAR BAY
13889         INLAND
5887    <1H OCEAN
19802         INLAND
2521         INLAND
2799         INLAND
9188         INLAND
Name: ocean_proximity, dtype: object

```

d. What is the ratio of bedrooms to total rooms in the data

```

# total no of rooms
total_rooms = data['total_rooms'].sum()
# total number of bedrooms
total_bedrooms = data['total_bedrooms'].sum()
#printing the ratio of bedrooms to total rooms
print(total_rooms//total_bedrooms)

```

Output

```
4.0
```

e. Determine the average price of a house for each type of ocean_proximity.

```

# average house price for each ocean_proximity type
data.groupby('ocean_proximity')['median_house_value'].median()

```

Output

```

ocean_proximity
<1H OCEAN    214850.0
INLAND       108500.0
ISLAND       414700.0
NEAR BAY     233800.0
NEAR OCEAN   229450.0
Name: median_house_value, dtype: float64

```

5. Write a program to perform the following tasks

a. Determine the outliers in each non-categorical column of Titanic Data and remove them.

```

# importing all the necessary libraries
import pandas as pd
import numpy as np

#we need to read the data
data = pd.read_csv("https://raw.githubusercontent.com/naveenjoshii/Intro-to-MachineLearning/master/Titanic/titanic.csv")
#print top 5 rows
print(data.head())

```

Output

	PassengerId	Survived	Pclass	...	Fare	Cabin	Embarked
0	1	0	3	...	7.2500	NaN	S
1	2	1	1	...	71.2833	C85	C
2	3	1	3	...	7.9250	NaN	S
3	4	1	1	...	53.1000	C123	S
4	5	0	3	...	8.0500	NaN	S

[5 rows x 12 columns]

function to calculate the lower and upperbound

def detect_outliers(data,threshold):

mean = np.mean(data)

std =np.std(data)

lb = max(mean - (threshold * std),min(data))

ub = min(mean + (threshold * std),max(data))

return lb,ub

df = data.copy()

lb,ub = detect_outliers(data["Fare"],4)

removing the rows which are greater than upperbound

df.drop(df[df.Fare > ub].index, inplace=True)

removing the rows which are less than lowerbound

df.drop(df[df.Fare < lb].index, inplace=True)

lb,ub = detect_outliers(data["Age"],5)

removing the rows which are greater than upperbound

df.drop(df[df.Age > ub].index, inplace=True)

removing the rows which are less than lowerbound

df.drop(df[df.Age < lb].index, inplace=True)

b. Determine missing values in each column of Titanic data. If missing values account for 30% of data, then remove the column.

#printing the missing value percentage for every column

df.isnull().mean() * 100

Output

PassengerId	0.000000
Survived	0.000000
Pclass	0.000000
Name	0.000000
Sex	0.000000
Age	20.113636
SibSp	0.000000
Parch	0.000000
Ticket	0.000000
Fare	0.000000

```
Cabin          77.954545
Embarked       0.227273
dtype: float64
```

```
# get all the column names in our dataset
df.columns
```

Output

```
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age',
      'SibSp',
      'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
      dtype='object')
```

```
# As we can see cabin column has more than 30% of missing values, so we have to drop that
column
df.drop(['Cabin'],inplace=True,axis=1)
```

```
# after removing the column cabin, printing the columns again. If you observe there is no
Cabin in the output
df.columns
```

Output

```
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age',
      'SibSp',
      'Parch', 'Ticket', 'Fare', 'Embarked'],
      dtype='object')
```

c. If missing values are less than 30% of entire data then create a new data frame
i. Missing values in numeric columns are filled with the mean of the corresponding column.

```
#printing the percentage of missing values in Age before handling
df['Age'].isnull().mean() * 100
```

Output

```
20.113636363636363
```

```
# Filling the missing values with the mean of respective column
df['Age']=df['Age'].fillna(df['Age'].mean())
```

```
#printing the percentage of missing values in Age after handling
df['Age'].isnull().mean() * 100
```

Output

```
0.0
```

ii. Missing values in categorical columns are filled with the most frequently occurring value.

```
#printing the percentage of missing values in Embarked before handling
df['Embarked'].isnull().mean() * 100
```

Output

```
0.22727272727272727
```

```
# filling with filled with the most frequently occurring value.
df["Embarked"].fillna(df['Embarked'].mode()[0],inplace=True)
```

```
#printing the percentage of missing values in Embarked after handling
df['Embarked'].isnull().mean() * 100
```

Output

```
0.0
```

6. Write a program to perform the following tasks

a. Determine the categorical columns in Titanic Dataset. Convert Columns with string data type to numerical data using encoding techniques.

```
#information about data
df.info()
```

Output

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 880 entries, 0 to 890
Data columns (total 11 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   PassengerId     880 non-null    int64
 1   Survived        880 non-null    int64
 2   Pclass          880 non-null    int64
 3   Name            880 non-null    object
 4   Sex             880 non-null    object
 5   Age            880 non-null    float64
 6   SibSp           880 non-null    int64
 7   Parch          880 non-null    int64
 8   Ticket         880 non-null    object
 9   Fare           880 non-null    float64
10   Embarked        880 non-null    object
dtypes: float64(2), int64(5), object(4)
memory usage: 122.5+ KB
```

```
print("each unique value and respective counts in Sex column\n",df['Sex'].value_counts())
```



```
#creating another data frame for Sex column
sex_df = pd.get_dummies(df['Sex'],drop_first=3)
sex_df.head()
```

Output

```
each unique value and respective counts in Sex column
male      572
female    308
Name: Sex, dtype: int64
```

```
male
0 1
1 0
2 0
3 0
4 1
```

```
print("each unique value and respective counts in Sex
column\n",df['Embarked'].value_counts())
# creating dummies for Embarked
embark_df = pd.get_dummies(df['Embarked'],drop_first=True)
embark_df.head()
```

Output

```
each unique value and respective counts in Sex column
S      642
C      161
Q       77
Name: Embarked, dtype: int64
```

```
Q S
0 0 1
1 0 0
2 0 1
3 0 1
4 0 1
```

```
old_data = df.copy()
# we need to drop the sex and embarked columns and replace them with the newly created
dummies data frames
# as Name and Ticket is not making any impact on the output label, we can drop them also
df.drop(['Sex','PassengerId','Embarked','Name','Ticket'],axis=1,inplace=True)
df.head()
```

Output

	Survived	Pclass	Age	SibSp	Parch	Fare
0	0	3	22.0	1	0	7.2500
1	1	1	38.0	1	0	71.2833

2	1	3	26.0	0	0	7.9250
3	1	1	35.0	1	0	53.1000
4	0	3	35.0	0	0	8.0500

```
# After dropping the Sex and Embarked columns, we are replacing them with out new data frames
data = pd.concat([df,sex_df,embark_df],axis=1)
```

b. Convert data in each numerical column so that it lies in the range [0,1]

```
# before scaling the data
data.head()
```

Output

	Survived	Pclass	Age	SibSp	Parch	Fare	male	Q	S
0	0	3	22.0	1	0	7.2500	1	0	1
1	1	1	38.0	1	0	71.2833	0	0	0
2	1	3	26.0	0	0	7.9250	0	0	1
3	1	1	35.0	1	0	53.1000	0	0	1
4	0	3	35.0	0	0	8.0500	1	0	1

```
# Scaling the data using minmax scaler so that values should be lies btw [0,1]
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
data[['Age','Pclass','Survived','SibSp','Parch','Fare','male','Q','S']] =
scaler.fit_transform(data[['Age','Pclass','Survived','SibSp','Parch','Fare','male','Q','S']])
```

```
# after scaling the data
data.head()
```

	Survived	Pclass	Age	SibSp	Parch	Fare	male	Q	S
0	0.0	1.0	0.271174	0.125	0.0	0.031865	1.0	0.0	1.0
1	1.0	0.0	0.472229	0.125	0.0	0.313299	0.0	0.0	0.0
2	1.0	1.0	0.321438	0.000	0.0	0.034831	0.0	0.0	1.0
3	1.0	0.0	0.434531	0.125	0.0	0.233381	0.0	0.0	1.0
4	0.0	1.0	0.434531	0.000	0.0	0.035381	1.0	0.0	1.0

7. Implement the following models on Titanic Dataset and determine the values of accuracy, precision, recall, f1 score and confusion matrix for the test data.

```
data.info()
```

Output

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 880 entries, 0 to 890
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Survived    880 non-null    float64
 1   Pclass      880 non-null    float64
 2   Age         880 non-null    float64
 3   SibSp       880 non-null    float64
 4   Parch       880 non-null    float64
 5   Fare        880 non-null    float64
 6   male        880 non-null    float64
 7   Q           880 non-null    float64
 8   S           880 non-null    float64
dtypes: float64(9)
memory usage: 108.8 KB

```

Split the Data

```

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(data.drop('Survived',axis=1),
                                                    data['Survived'], test_size=0.30,
                                                    random_state=101)

```

a. Logistic Regression

```

from sklearn.linear_model import LogisticRegression

# Build the Model.
logmodel = LogisticRegression()
logmodel.fit(X_train,y_train)

```

Output

```

LogisticRegression(C=1.0, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001,
verbose=0,
                    warm_start=False)

```

```

print("Predicting the model on the test set")
predicted = logmodel.predict(X_test)

```

Output

```

Predicting the model on the test set

```

```

print("predicted result !")

```

predicted

Output

predicted result !

```
array([1., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 1., 0., 0., 0., 0., 1.,
       0., 1., 1., 1., 0., 0., 1., 0., 1., 0., 0., 0., 1., 1., 0., 0., 1.,
       0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 1., 1., 0., 0., 0., 0.,
       0., 1., 0., 0., 1., 1., 0., 0., 0., 1., 1., 0., 0., 0., 1., 0., 0.,
       0., 1., 1., 0., 1., 0., 0., 0., 0., 0., 1., 0., 0., 1., 0., 1., 0.,
       1., 0., 1., 0., 0., 1., 0., 0., 0., 1., 0., 1., 1., 1., 0., 0., 1.,
       0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 1., 0., 0.,
       0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 1., 1., 1., 0., 1., 0.,
       0., 1., 0., 1., 0., 0., 0., 0., 1., 0., 0., 0., 0., 1., 0., 0., 0.,
       1., 0., 0., 1., 1., 0., 1., 0., 1., 0., 0., 0., 1., 0., 1., 0., 0.,
       0., 1., 0., 0., 1., 1., 1., 0., 0., 0., 0., 1., 0., 0., 0., 0., 1.,
       0., 0., 0., 0., 1., 0., 1., 0., 1., 1., 1., 1., 0., 0., 1., 1., 0.,
       0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 1., 0., 1., 0., 1., 0., 1.,
       1., 0., 0., 1., 0., 1., 1., 1., 1., 1., 0., 1., 0., 1., 1., 1., 1.,
       1., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0.,
       1., 1., 0., 1., 0., 0., 1., 0., 0.] )
```

#confusion matrix

```
from sklearn.metrics import confusion_matrix, classification_report
print(confusion_matrix(y_test, predicted))
```

Output

```
[[144  24]
 [ 28  68]]
```

Precision Score

```
from sklearn.metrics import precision_score
print("Precision Score",precision_score(y_test,predicted))
```

Output

Precision Score 0.7391304347826086

Recall Score

```
from sklearn.metrics import recall_score
print("recall score",recall_score(y_test,predicted))
```

Output

recall score 0.7083333333333334

F1 Score

```
from sklearn.metrics import f1_score
print("f1 score",f1_score(y_test,predicted))
```

Output

```
f1 score 0.723404255319149
```

Classification report

```
from sklearn.metrics import classification_report
print(classification_report(y_test,predicted))
```

Output

	precision	recall	f1-score	support
0.0	0.84	0.86	0.85	168
1.0	0.74	0.71	0.72	96
accuracy			0.80	264
macro avg	0.79	0.78	0.79	264
weighted avg	0.80	0.80	0.80	264

metrics are used to find accuracy or error

```
from sklearn import metrics
```

using metrics module for accuracy calculation

```
print("ACCURACY of Logistic Regression Model: ", metrics.accuracy_score(y_test,
predicted))
```

Output

```
ACCURACY of Logistic Regression Model: 0.803030303030303
```

b. Random Forest Classifier

importing random forest classifier from assemble module

```
from sklearn.ensemble import RandomForestClassifier
```

creating a RF classifier

```
clf = RandomForestClassifier(n_estimators = 100)
```

Training the model on the training dataset

fit function is used to train the model using the training sets as parameters

```
clf.fit(X_train, y_train)
```

performing predictions on the test dataset

```
y_pred = clf.predict(X_test)
```

#confusion matrix

```
from sklearn.metrics import confusion_matrix, classification_report
print(confusion_matrix(y_test, y_pred))
```

Output

```
[[140  28]
 [ 20  76]]
```

Precision Score

```
from sklearn.metrics import precision_score
print("Precision Score",precision_score(y_test,y_pred))
```

Output

```
Precision Score 0.7307692307692307
```

Recall Score

```
from sklearn.metrics import recall_score
print("recall score",recall_score(y_test,y_pred))
```

Output

```
recall score 0.7916666666666666
```

F1 Score

```
from sklearn.metrics import f1_score
print("f1 score",f1_score(y_test,y_pred))
```

Output

```
f1 score 0.76
```

Classification report

```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

Output

	precision	recall	f1-score	support
0.0	0.88	0.83	0.85	168
1.0	0.73	0.79	0.76	96
accuracy			0.82	264
macro avg	0.80	0.81	0.81	264
weighted avg	0.82	0.82	0.82	264

metrics are used to find accuracy or error

from sklearn import metrics

using metrics module for accuracy calculation

```
print("ACCURACY of Random Forest Classifier Model: ", metrics.accuracy_score(y_test,
y_pred))
```

Output

```
ACCURACY of Random Forest Classifier Model: 0.8181818181818182
```

8. Implement the following models on the California House Pricing Dataset and determine the values of R2 score, the area under roc curve and root mean squared error for the test set.

a. Linear Regression with Polynomial Features

b. Random Forest Regressor

Preparing the data

```
# checking for null values
data.isnull().mean() * 100
```

Output

```
longitude          0.000000
latitude           0.000000
housing_median_age  0.000000
total_rooms         0.000000
total_bedrooms      1.002907
population          0.000000
households          0.000000
median_income       0.000000
median_house_value  0.000000
ocean_proximity     0.000000
diff_income_and_house_value  0.000000
dtype: float64
```

```
# handling null values in total_bedrooms with the most frequent value in respective column
data["total_bedrooms"].fillna(data['total_bedrooms'].mode()[0],inplace=True)
```

```
#checking the null values handled or not
data["total_bedrooms"].isnull().mean() * 100
```

Output

```
0.0
```

```
data.info()
```

Output

```
data.info()
data.info()
<class 'pandas.core.frame.DataFrame'>
```

Int64Index: 20640 entries, 4861 to 9188

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	longitude	20640 non-null	float64
1	latitude	20640 non-null	float64
2	housing_median_age	20640 non-null	float64
3	total_rooms	20640 non-null	float64
4	total_bedrooms	20640 non-null	float64
5	population	20640 non-null	float64
6	households	20640 non-null	float64
7	median_income	20640 non-null	float64
8	median_house_value	20640 non-null	float64
9	ocean_proximity	20640 non-null	object
10	diff_income_and_house_value	20640 non-null	float64

dtypes: float64(10), object(1)

memory usage: 1.9+ MB

data['ocean_proximity'].unique()

Output

```
array(['<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'NEAR BAY', 'ISLAND'],  
      dtype=object)
```

#we need to convert categorical values by label encoding

there are more than two categories, we have to use onehot encoding

data['ocean_proximity'].value_counts()

ocean_prox_df = pd.get_dummies(data['ocean_proximity'],drop_first=True)

ocean_prox_df.head()

Output

	INLAN D	ISLAN D	NEAR BAY	NEAR OCEAN
4861	0	0	0	0
6688	1	0	0	0
16642	0	0	0	1
15661	0	0	1	0
15652	0	0	1	0

old_data = data.copy()


```
data.drop(['ocean_proximity','longitude','latitude','diff_income_and_house_value'],axis=1,inplace=True)
data.head()
```

Output

	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
4861	29.0	515.0	229.0	2690.0	217.0	0.4999	500001.0
6688	28.0	238.0	58.0	142.0	31.0	0.4999	500001.0
16642	19.0	1540.0	715.0	1799.0	635.0	0.7025	500001.0
15661	27.0	1728.0	884.0	1211.0	752.0	0.8543	500001.0
15652	52.0	3260.0	1535.0	3260.0	1457.0	0.9000	500001.0

```
data = pd.concat([data,ocean_prox_df],axis=1)
```

```
data.head()
```

Output

	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
4861	29.0	515.0	229.0	2690.0	217.0	0.4999	500001.0	0	0	0	0
6688	28.0	238.0	58.0	142.0	31.0	0.4999	500001.0	1	0	0	0
16642	19.0	1540.0	715.0	1799.0	635.0	0.7025	500001.0	0	0	0	1
15661	27.0	1728.0	884.0	1211.0	752.0	0.8543	500001.0	0	0	1	0
15652	52.0	3260.0	1535.0	3260.0	1457.0	0.9000	500001.0	0	0	1	0

Split the data

```
from sklearn.model_selection import train_test_split
# split the data for training and testing
```

```
X_train, X_test, y_train, y_test = train_test_split(data.drop('median_house_value',axis=1),
                                                    data['median_house_value'], test_size=0.30,
                                                    random_state=101)
```

a. Linear Regression with Polynomial Features

```
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
```

```
#model initialization
model = LinearRegression()

# initializing polynomial featurers
poly = PolynomialFeatures(degree=3)
#converting features into polyfeatures
X_ = poly.fit_transform(X_train)
Y_ = poly.fit_transform(y_train.values.reshape(-1,1))
# training the model
model.fit(X_,Y_)
```

Output

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

```
#preparing test data for predictions
testX = poly.fit_transform(X_test)
# predicting the output for test data
predicted = model.predict(testX)
```

```
# expected output for test data
expected = poly.fit_transform(y_test.values.reshape(-1,1))
```

```
from sklearn.metrics import r2_score
r2 = r2_score(expected, predicted)
print('r2 score is', r2)
```

Output

```
r2 score is 0.590661764648472
```

```
# example of calculate the root mean squared error
from sklearn.metrics import mean_squared_error
# calculate errors
errors = mean_squared_error(expected, predicted, squared=False)
# report error
print("root mean square error is :",errors)
```

Output

```
root mean square error is : 1.1921996852169048e+16
```

b. Random Forest Regressor

```
# Fitting Random Forest Regression to the dataset
# import the regressor
from sklearn.ensemble import RandomForestRegressor

# create regressor object
regressor = RandomForestRegressor(n_estimators = 100, random_state = 101)

# fit the regressor with x and y data
regressor.fit(X_train, y_train)
```

Output

```
RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                        max_depth=None, max_features='auto',
                        max_leaf_nodes=None,
                        max_samples=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2,
                        min_weight_fraction_leaf=0.0,
                        n_estimators=100, n_jobs=None, oob_score=False,
                        random_state=101, verbose=0, warm_start=False)
```

```
# test the output by changing values
predicted = regressor.predict(X_test)
```

```
expected = y_test
```

```
from sklearn.metrics import r2_score
r2 = r2_score(expected, predicted)
print('r2 score is', r2)
```

Output

```
r2 score is 0.7091234171276952
```

```
# example of calculate the root mean squared error
from sklearn.metrics import mean_squared_error
# calculate errors
errors = mean_squared_error(expected, predicted, squared=False)
# report error
print("root mean square error is :", errors)
```

Output

```
root mean square error is : 62360.02542136252
```

1. Implement a single neural network and test for different logic gates.

#Or gate

import numpy as np

def unitStep(v):

if v >= 0:

return 1

else:

return 0

def perceptronModel(x, w, b):

v = np.dot(w, x) + b

y = unitStep(v)

return y

OR Logic Function

w1 = 1, w2 = 1, b = -0.5

def OR_logicFunction(x):

w = np.array([1, 1])

b = -0.5

return perceptronModel(x, w, b)

testing the Perceptron Model

test1 = np.array([0, 1])

test2 = np.array([1, 1])

test3 = np.array([0, 0])

test4 = np.array([1, 0])

print("OR({}, {}) = {}".format(0, 1, OR_logicFunction(test1)))

print("OR({}, {}) = {}".format(1, 1, OR_logicFunction(test2)))

print("OR({}, {}) = {}".format(0, 0, OR_logicFunction(test3)))

print("OR({}, {}) = {}".format(1, 0, OR_logicFunction(test4)))

Output

```
OR(0, 1) = 1
OR(1, 1) = 1
OR(0, 0) = 0
OR(1, 0) = 1
```

And gate

```
import numpy as np
# define Unit Step Function
def unitStep(v):
    if v >= 0:
        return 1
    else:
        return 0
```

design Perceptron Model

```
def perceptronModel(x, w, b):
    v = np.dot(w, x) + b
    y = unitStep(v)
    return y
```

AND Logic Function

```
# w1 = 1, w2 = 1, b = -1.5
def AND_logicFunction(x):
    w = np.array([1, 1])
    b = -1.5
    return perceptronModel(x, w, b)
```

testing the Perceptron Model

```
test1 = np.array([0, 1])
test2 = np.array([1, 1])
test3 = np.array([0, 0])
test4 = np.array([1, 0])
```

```
print("AND({}, {}) = {}".format(0, 1, AND_logicFunction(test1)))
print("AND({}, {}) = {}".format(1, 1, AND_logicFunction(test2)))
print("AND({}, {}) = {}".format(0, 0, AND_logicFunction(test3)))
print("AND({}, {}) = {}".format(1, 0, AND_logicFunction(test4)))
```

Output

```
AND(0, 1) = 0
AND(1, 1) = 1
AND(0, 0) = 0
AND(1, 0) = 0
```

2. Write a program to train and test a Convolutional Neural Network to determine the number, given an image of a handwritten digit. Determine the training and validation accuracies of your model. (Train your model for 5 epochs).

```

from keras.datasets import mnist

# loading the dataset
(X_train, y_train), (X_test, y_test) = mnist.load_data()
# let's print the shape of the dataset

```

Output

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11493376/11490434 [=====] - 0s 0us/step
11501568/11490434 [=====] - 0s 0us/step

```

```

print("X_train shape", X_train.shape)
print("y_train shape", y_train.shape)
print("X_test shape", X_test.shape)
print("y_test shape", y_test.shape)

```

Output

```

X_train shape (60000, 28, 28)
y_train shape (60000,)
X_test shape (10000, 28, 28)
y_test shape (10000,)

```

```

# keras imports for the dataset and building our neural network
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Conv2D, MaxPool2D
from keras.utils import np_utils

```

```

# Flattening the images from the 28x28 pixels to 1D 784 pixels
X_train = X_train.reshape(60000, 784)
X_test = X_test.reshape(10000, 784)
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')

```

```

# normalizing the data to help with the training
X_train /= 255
X_test /= 255

```

```

# one-hot encoding using keras' numpy-related utilities
n_classes = 10
print("Shape before one-hot encoding: ", y_train.shape)
Y_train = np_utils.to_categorical(y_train, n_classes)
Y_test = np_utils.to_categorical(y_test, n_classes)
print("Shape after one-hot encoding: ", Y_train.shape)

```

```
# building a linear stack of layers with the sequential model
model = Sequential()
# hidden layer
model.add(Dense(100, input_shape=(784,), activation='relu'))
# output layer
model.add(Dense(10, activation='softmax'))

# looking at the model summary
model.summary()
# compiling the sequential model
model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')
# training the model for 10 epochs
model.fit(X_train, Y_train, batch_size=128, epochs=10, validation_data=(X_test, Y_test))
```

```
Shape before one-hot encoding: (60000,)
Shape after one-hot encoding: (60000, 10)
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 100)	78500
dense_1 (Dense)	(None, 10)	1010

```
=====
Total params: 79,510
Trainable params: 79,510
Non-trainable params: 0
```

```
Epoch 1/10
469/469 [=====] - 3s 5ms/step - loss: 0.3805 -
accuracy: 0.8950 - val_loss: 0.2060 - val_accuracy: 0.9409
Epoch 2/10
469/469 [=====] - 2s 5ms/step - loss: 0.1812 -
accuracy: 0.9477 - val_loss: 0.1493 - val_accuracy: 0.9566
Epoch 3/10
469/469 [=====] - 2s 5ms/step - loss: 0.1334 -
accuracy: 0.9613 - val_loss: 0.1223 - val_accuracy: 0.9644
Epoch 4/10
469/469 [=====] - 2s 5ms/step - loss: 0.1055 -
accuracy: 0.9699 - val_loss: 0.1059 - val_accuracy: 0.9693
Epoch 5/10
469/469 [=====] - 2s 5ms/step - loss: 0.0863 -
accuracy: 0.9753 - val_loss: 0.1025 - val_accuracy: 0.9697
Epoch 6/10
469/469 [=====] - 2s 4ms/step - loss: 0.0718 -
accuracy: 0.9796 - val_loss: 0.0951 - val_accuracy: 0.9721
Epoch 7/10
469/469 [=====] - 2s 4ms/step - loss: 0.0615 -
accuracy: 0.9822 - val_loss: 0.0865 - val_accuracy: 0.9735
Epoch 8/10
469/469 [=====] - 2s 5ms/step - loss: 0.0535 -
accuracy: 0.9851 - val_loss: 0.0800 - val_accuracy: 0.9761
Epoch 9/10
469/469 [=====] - 2s 4ms/step - loss: 0.0457 -
accuracy: 0.9868 - val_loss: 0.0829 - val_accuracy: 0.9754
```

```
Epoch 10/10
469/469 [=====] - 2s 4ms/step - loss: 0.0391 -
accuracy: 0.9888 - val_loss: 0.0784 - val_accuracy: 0.9757
```

Output

```
<keras.callbacks.History at 0x7f6bd453df10>
```

```
# keras imports for the dataset and building our neural network
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Conv2D, MaxPool2D, Flatten
from keras.utils import np_utils
```

```
# to calculate accuracy
from sklearn.metrics import accuracy_score
```

```
# loading the dataset
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
# building the input vector from the 28x28 pixels
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
```

```
# normalizing the data to help with the training
X_train /= 255
X_test /= 255
```

```
# one-hot encoding using keras' numpy-related utilities
n_classes = 10
print("Shape before one-hot encoding: ", y_train.shape)
Y_train = np_utils.to_categorical(y_train, n_classes)
Y_test = np_utils.to_categorical(y_test, n_classes)
print("Shape after one-hot encoding: ", Y_train.shape)
```

```
# building a linear stack of layers with the sequential model
model = Sequential()
# convolutional layer
model.add(Conv2D(25, kernel_size=(3,3), strides=(1,1), padding='valid', activation='relu',
input_shape=(28,28,1)))
model.add(MaxPool2D(pool_size=(1,1)))
# flatten output of conv
model.add(Flatten())
# hidden layer
model.add(Dense(100, activation='relu'))
# output layer
model.add(Dense(10, activation='softmax'))
```

```
# compiling the sequential model
```



```
model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

# training the model for 10 epochs
model.fit(X_train, Y_train, batch_size=128, epochs=10, validation_data=(X_test, Y_test))
```

```
Shape before one-hot encoding: (60000,)
Shape after one-hot encoding: (60000, 10)
Epoch 1/10
469/469 [=====] - 41s 86ms/step - loss: 0.2190 -
accuracy: 0.9367 - val_loss: 0.0841 - val_accuracy: 0.9768
Epoch 2/10
469/469 [=====] - 42s 90ms/step - loss: 0.0659 -
accuracy: 0.9804 - val_loss: 0.0538 - val_accuracy: 0.9820
Epoch 3/10
469/469 [=====] - 40s 84ms/step - loss: 0.0376 -
accuracy: 0.9891 - val_loss: 0.0527 - val_accuracy: 0.9827
Epoch 4/10
469/469 [=====] - 40s 86ms/step - loss: 0.0243 -
accuracy: 0.9926 - val_loss: 0.0563 - val_accuracy: 0.9806
Epoch 5/10
469/469 [=====] - 40s 84ms/step - loss: 0.0152 -
accuracy: 0.9956 - val_loss: 0.0598 - val_accuracy: 0.9834
Epoch 6/10
469/469 [=====] - 40s 85ms/step - loss: 0.0104 -
accuracy: 0.9968 - val_loss: 0.0579 - val_accuracy: 0.9826
Epoch 7/10
469/469 [=====] - 40s 85ms/step - loss: 0.0070 -
accuracy: 0.9983 - val_loss: 0.0661 - val_accuracy: 0.9828
Epoch 8/10
469/469 [=====] - 40s 85ms/step - loss: 0.0056 -
accuracy: 0.9983 - val_loss: 0.0542 - val_accuracy: 0.9842
Epoch 9/10
469/469 [=====] - 40s 85ms/step - loss: 0.0046 -
accuracy: 0.9989 - val_loss: 0.0674 - val_accuracy: 0.9833
Epoch 10/10
469/469 [=====] - 40s 85ms/step - loss: 0.0052 -
accuracy: 0.9985 - val_loss: 0.0720 - val_accuracy: 0.9818
```

Output

```
<keras.callbacks.History at 0x7f6bcfde47d0>
```